

A perfectly secure symmetric encryption scheme: ONE-TIME PAD

This encryption scheme achieves information-theoretic security.

Definition 1 (Symmetric encryption). Let \mathcal{K} be a key space, \mathcal{P} be a plain-text space and let \mathcal{C} be a ciphertext space. These three spaces are finite spaces.

A *symmetric encryption* scheme over $(\mathcal{K}, \mathcal{P}, \mathcal{C})$ is a tuple of three algorithms (KeyGen, Enc, Dec) :

- ▷ KeyGen provides a sample k of \mathcal{K} ;
- ▷ $\text{Enc} : \mathcal{K} \times \mathcal{P} \rightarrow \mathcal{C}$;
- ▷ $\text{Dec} : \mathcal{K} \times \mathcal{C} \rightarrow \mathcal{P}$.

Without loss of generality, we will assume that $\text{im Enc} = \mathcal{C}$. We want to ensure **Correctness**: for any key $k \in \mathcal{K}$ and message $m \in \mathcal{P}$, we have that:

$$\text{Dec}(k, \text{Enc}(k, m)) = m.$$

The elements m and k are independent random variables and all the elements in \mathcal{K} and \mathcal{P} have non-zero probability.

Remark 1. The algorithm Enc could (and should¹) be probabilistic. However, the algorithm Dec is deterministic.

So far, we did not talk about efficiency of these algorithms.

Definition 2 (Shannon, 1949). A symmetric encryption scheme is said to have *perfect security* whenever, for any \bar{m} and any \bar{c} ,

$$\Pr_{k,m}[m = \bar{m} \mid \text{Enc}_k(m) = \bar{c}] = \Pr_m[m = \bar{m}].$$

The intuition is that knowing the encrypted message tells me *nothing* about the message.

Lemma 1 (Shannon). Given a symmetric encryption scheme (KeyGen, Enc, Dec) has perfect security then $|\mathcal{K}| \geq |\mathcal{P}|$.

Proof. Let $\bar{c} \in \mathcal{C}$ and define

$$\mathcal{S} := \{\bar{m} \in \mathcal{P} \mid \exists \bar{k} \in \mathcal{K}, \bar{m} = \text{Dec}(\bar{k}, \bar{c})\}.$$

Let $N := |\mathcal{S}|$. We have that $N \leq |\mathcal{K}|$ as Dec is deterministic. We also have that $N \leq |\mathcal{P}|$ as $\mathcal{S} \subseteq \mathcal{P}$. Finally, assume $N < |\mathcal{P}|$. This means, there exists $\bar{m} \in \mathcal{P}$ such that $\bar{m} \notin \mathcal{S}$. Then,

$$\Pr[m = \bar{m} \mid \text{Enc}_k(m) = \bar{c}] = 0,$$

but by assumption, $\Pr[m = \bar{m}] \neq 0$. So this is not a perfectly secure scheme. We can conclude that

$$N = |\mathcal{P}| \leq |\mathcal{K}|.$$

□

¹If the algorithm is deterministic, if we see two identical ciphers we know that the messages are identical, and this can be seen as a vulnerability of this protocol.

Example 1 (One-Time PAD). Let $\mathcal{K} = \mathcal{C} = \mathcal{P} = \{0, 1\}^\ell$. Here are the algorithms used:

- ▷ KeyGen samples from $\mathcal{U}(\{0, 1\}^\ell)$.
- ▷ Enc(k, m) we compute the XOR $c = m \oplus k$.
- ▷ Dec(k, m) we compute the XOR $m = c \oplus k$.

Theorem 1. The One-Time PAD is a perfectly-secure symmetric encryption.

Proof. Correctness. We have that

$$\text{Dec}(k, \text{Enc}(k, m)) = k \oplus k \oplus m = m.$$

Security. We have, by independence of m and k we have that

$$\begin{aligned} \Pr[m = \bar{m} \mid \text{Enc}(k, m) = \bar{c}] &= \Pr[m = \bar{m} \mid k \oplus m = \bar{c}] \\ &= \Pr[m = \bar{m}]. \end{aligned}$$

□

Remark 2. This example is not practical:

- ▷ keys need to be larger than the message;
- ▷ you cannot encrypt twice: for example, $c_1 = m_1 \oplus k$ and $c_2 = m_2 \oplus k$, then we have $c_1 \oplus c_2 = m_1 \oplus m_2$.

This last part is why that protocol is called a *One-Time secure encryption*.

We want to be able to encrypt arbitrarily long messages! We will have to make a trade-off and we choose to not care about *perfect* security. Why? In real life, we don't care about proving that something is proven to be absolutely infeasible, we only want to believe it is

infeasible in practice.

Computational complexity is sufficient in practice.

Let us be more precise.

Definition 3. Let \mathcal{D}_0 and \mathcal{D}_1 be two distributions over $\{0, 1\}^n$.

An algorithm $\mathcal{A} : \{0, 1\}^n \rightarrow \{0, 1\}$ is called a *distinguisher* between \mathcal{D}_0 and \mathcal{D}_1 . We define its *distinguishing advantage* as:

$$\text{Adv}_{\mathcal{A}} := \left| \underbrace{\Pr_{x \leftarrow \mathcal{D}_1} [\mathcal{A}(X) = 1]}_{\text{probability of being right}} - \underbrace{\Pr_{x \leftarrow \mathcal{D}_0} [\mathcal{A}(X) = 1]}_{\text{probability of being mistaken}} \right|.$$

We say that \mathcal{D}_0 and \mathcal{D}_1 are *computationally indistinguishable* if for any efficient distinguisher \mathcal{A} its advantage $\text{Adv}_{\mathcal{A}}$ is small.

This definition is not very formal yet, we have not defined “efficient” and “small.” This can be formalized by introducing a parameter $\lambda \in \mathbb{N}$ called the *security parameter*.

Definition 4. Let $(\mathcal{D}_{0,\lambda})_{\lambda \in \mathbb{N}}$ and $(\mathcal{D}_{1,\lambda})_{\lambda \in \mathbb{N}}$ be two distributions over $\{0, 1\}^{n(\lambda)}$ for a non-decreasing polynomial $n(\lambda)$. The value of $\lambda \in \mathbb{N}$ is called the *security parameter*.

An algorithm $\mathcal{A} : \{0, 1\}^{n(\lambda)} \rightarrow \{0, 1\}$ is called a *distinguisher* between the distributions $\mathcal{D}_{0,\lambda}$ and $\mathcal{D}_{1,\lambda}$. We define its *distinguishing advantage* as:

$$\text{Adv}_{\mathcal{A}}(\lambda) := \left| \underbrace{\Pr_{x \leftarrow \mathcal{D}_{1,\lambda}} [\mathcal{A}(X) = 1]}_{\text{probability of being right}} - \underbrace{\Pr_{x \leftarrow \mathcal{D}_{0,\lambda}} [\mathcal{A}(X) = 1]}_{\text{probability of being mistaken}} \right|.$$

We say that $\mathcal{D}_{0,\lambda}$ and $\mathcal{D}_{1,\lambda}$ are *computationally indistinguishable* if for any distinguisher \mathcal{A} running in $O(\lambda^c)$ for some $c > 0$ ² its advantage $\text{Adv}_{\mathcal{A}}$ is a $o(1/\lambda^c)$ for some $c > 0$.³

Our goal now is to extend the One-Time PAD to messages m larger than the key k . We want to construct some function G that takes as input the key $k \in \{0, 1\}^n$ and expand it to a string $G(k) \in \{0, 1\}^\ell$ for some $\ell > n$ that is computationally hard to distinguish from a uniform random string. This is called a *PRG* or *pseudo-random generator*.

Definition 5. A *pseudo-random generator* is a pair of poly-time algorithms (Setup, G) such that:

- ▷ Setup is an algorithm that takes as input a security parameter λ (taken as a string 1^λ of length λ , *i.e.* we write λ in unary) and returns a public parameter;
- ▷ $G_\lambda : \{0, 1\}^{n(\lambda)} \rightarrow \{0, 1\}^{\ell(\lambda)}$ is an algorithm which takes a string k of length $n(\lambda)$ and return a string $G(k)$ of length $\ell(\lambda)$ with $\ell(\lambda) > n(\lambda)$.

such that

- ▷ G is deterministic;
- ▷ $\ell(\lambda) > n(\lambda)$ (we say that it is *expanding*)
- ▷ the distributions $\{\mathcal{U}(\{0, 1\}^{\ell(\lambda)})\}_{\lambda \in \mathbb{N}}$ and $\{G(\mathcal{U}(\{0, 1\}^{n(\lambda)}))\}_{\lambda \in \mathbb{N}}$ are computationally indistinguishable (we call it *pseudo-randomness*).

Another way of defining a pseudo-random generator is with *unpredictability* instead of *pseudo-randomness*.

Definition 6. This is the same definition as before but replacing pseudo-randomness with *unpredictability*.

A PRG (Setup, G) is *unpredictable* if, for any index $i \in \{0, \dots, \ell(\lambda)\}$

²This means it is polynomial in λ , which we will write $\text{poly}(\lambda)$

³This means it is negligible in terms of λ , which we will write $\text{negl}(\lambda)$.

and any efficient adversary $\mathcal{A} : \{0, 1\}^n \rightarrow \{0, 1\}$, we have that:

$$\left| \Pr_{k \leftarrow \mathcal{U}(\{0,1\}^{n(\lambda)})} [\mathcal{A}(G(k)_{|i}) = G(k)_{i+1}] - \frac{1}{2} \right| = \text{negl}(\lambda).$$

We can now prove that the two definitions are equivalent.

Theorem 2. The two definitions of a PRG are equivalent.

Proof. To simplify, we will remove the security parameter from the notations.

On one side, assume we have a predictor $\mathcal{A} : \{0, 1\}^i \rightarrow \{0, 1\}$ that succeeds in guessing $G(k)_{i+1}$ with non-negligible probability. We then construct a distinguisher \mathcal{B} against pseudo-randomness as \mathcal{B} receive a sample x from either $\mathcal{D}_0 = \mathcal{U}(\{0, 1\}^\ell)$ or $\mathcal{D}_1 = G(\mathcal{U}(\{0, 1\}^n))$: algorithm \mathcal{B} runs \mathcal{A} on input $x_{|i}$ and checks if $\mathcal{A}(x_{|i}) \stackrel{?}{=} x_{i+1}$. In that case, \mathcal{B} will return 1; otherwise it returns 0. What is the advantage of \mathcal{B} ?

$$\begin{aligned} \text{Adv}_{\mathcal{B}} &= \left| \Pr_{x \leftarrow \mathcal{D}_1} [\mathcal{B}(x) = 1] - \overbrace{\Pr_{x \leftarrow \mathcal{D}_0} [\mathcal{B}(x) = 1]}^{1/2} \right| \\ &= \left| \Pr_{x \leftarrow \mathcal{D}_1} [\mathcal{A}(x_{|i}) = x_{i+1}] - \frac{1}{2} \right|. \end{aligned}$$

This is the definition of the predictability advantage of \mathcal{A} (which is non-negligible by assumption).

Next, we will use a technique called an *Hybrid Argument* (due to Yao in '82). Assume we have a distinguisher \mathcal{A} such that

$$\text{Adv}_{\mathcal{A}} = \left| \Pr_{x \leftarrow \mathcal{D}_1} [\mathcal{A}(x) = 1] - \Pr_{x \leftarrow \mathcal{D}_0} [A(x) = 1] \right|$$

is non-negligible, say $\text{Adv}_{\mathcal{A}} \geq \varepsilon$. We then define $\ell+1$ distributions

$(\mathcal{D}_i)_{i=0,\dots,\ell}$ as

$$\mathcal{D}_i := \left\{ x \in \{0, 1\}^\ell \mid \begin{array}{l} x_{|i} = G(k)_{|i} \text{ for } k \leftarrow \mathcal{U}(\{0, 1\}^n) \\ x_{|i+1,\dots,\ell} \leftarrow \mathcal{U}(\{0, 1\}^{\ell-i}) \end{array} \right\}.$$

We then have, by all the terms cancelling (this is a telescoping sum), that:

$$\begin{aligned} \varepsilon \leq \text{Adv}_{\mathcal{A}}(\mathcal{D}_0, \mathcal{D}_n) &= \left| \sum_{i=0}^{\ell} \left(\Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] - \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] \right) \right| \\ &\leq \sum_{i=0}^{\ell} \left| \Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] - \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] \right| \\ &\leq \sum_{i=0}^{\ell} \text{Adv}_{\mathcal{A}}(\mathcal{D}_i, \mathcal{D}_{i+1}). \end{aligned}$$

By the pigeonhole principle, we have that there exists an $i \in \{0, \dots, \ell\}$, such that

$$\left| \Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] - \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] \right| \geq \frac{\varepsilon}{\ell + 1}.$$

As ε is non-negligible and $\ell + 1$ being polynomial in λ , we have that $\varepsilon/(\ell + 1)$ is non-negligible. How to turn this into a predictor for i ? Let us define \mathcal{B}_i as a predictor which is given $G(k)_{|i}$ and supposed to predict $G(k)_{i+1}$. Algorithm \mathcal{B}_i will compute $x \in \{0, 1\}^\ell$ with $x \leftarrow G(k)_{|i} || y$ where $y \leftarrow \mathcal{U}(\{0, 1\}^{\ell-i})$. Then \mathcal{B}_i runs algorithm \mathcal{A} on input x , and \mathcal{A} returns a bit $b \in \{0, 1\}$ and \mathcal{B}_i outputs a prediction x_{i+1} for $G(k)_{i+1}$ if $b = 1$ and $1 - x_{i+1}$

otherwise. What is the prediction advantage of \mathcal{B}_i ?

$$\begin{aligned}
 & \Pr[\mathcal{B}_i(G(k)|_i) = G(k)_{i+1}] \\
 &= \Pr \left[\begin{array}{c} \mathcal{A}(x) = 0 \wedge x_{i+1} = 1 - G(k)_{i+1} \\ \vee \\ \mathcal{A}(x) = 1 \wedge x_{i+1} = G(k)_{i+1} \end{array} \right] \\
 &= \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 0 \wedge x_{i+1} = 1 - G(k)_{i+1}] \\
 &\quad + \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1 \wedge x_{i+1} = G(k)_{i+1}] \\
 &= \frac{1}{2} \Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 0] + \frac{1}{2} \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] \\
 &= \frac{1}{2} \left(\Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] + 1 - \Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] \right)
 \end{aligned}$$

where we write $\bar{\mathcal{D}}_{i+1}$ is the “flipped” of \mathcal{D}_{i+1} . We have that:

$$\begin{aligned}
 & \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] \\
 &= \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1 \wedge x_{i+1} = G(k)_{i+1}] \\
 &\quad + \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1 \wedge x_{i+1} = 1 - G(k)_{i+1}] \\
 &= \frac{1}{2} \left(\Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] + \Pr_{x \leftarrow \bar{\mathcal{D}}_{i+1}} [\mathcal{A}(x) = 1] \right)
 \end{aligned}$$

thus

$$\Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] = 2 \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] - \Pr_{x \leftarrow \bar{\mathcal{D}}_{i+1}} [\mathcal{A}(x) = 1].$$

Hence,

$$\begin{aligned}
 & \Pr[\mathcal{B}_i(G(k)|_i) - G(k)_{i+1}] = \\
 & \frac{1}{2} \Pr_{x \leftarrow \mathcal{D}_{i+1}} [\mathcal{A}(x) = 1] + 1 - 2 \Pr_{x \leftarrow \mathcal{D}_i} [\mathcal{A}(x) = 1] + \Pr_{x \leftarrow \bar{\mathcal{D}}_{i+1}} [\mathcal{A}(x) = 1].
 \end{aligned}$$

Finally, we can conclude that:

$$\text{Adv}_{\mathcal{A}}(\mathcal{D}_i, \mathcal{D}_{i+1}) = \left| \Pr[\mathcal{B}_i(G(k)_{|i}) = G(k)_{i+1}] - \frac{1}{2} \right| \geq \frac{\varepsilon}{n}.$$

□